

Part III

Counter measures

HOW DO WE STOP THE ATTACKS?

- The best defense is proper bounds checking
- but there are many C/C++ programmers and some are bound to forget



➔ Are there any *system* defenses that can help?

HOW DO WE STOP THE ATTACKS?

- A variety of tricks in combination


NX bit



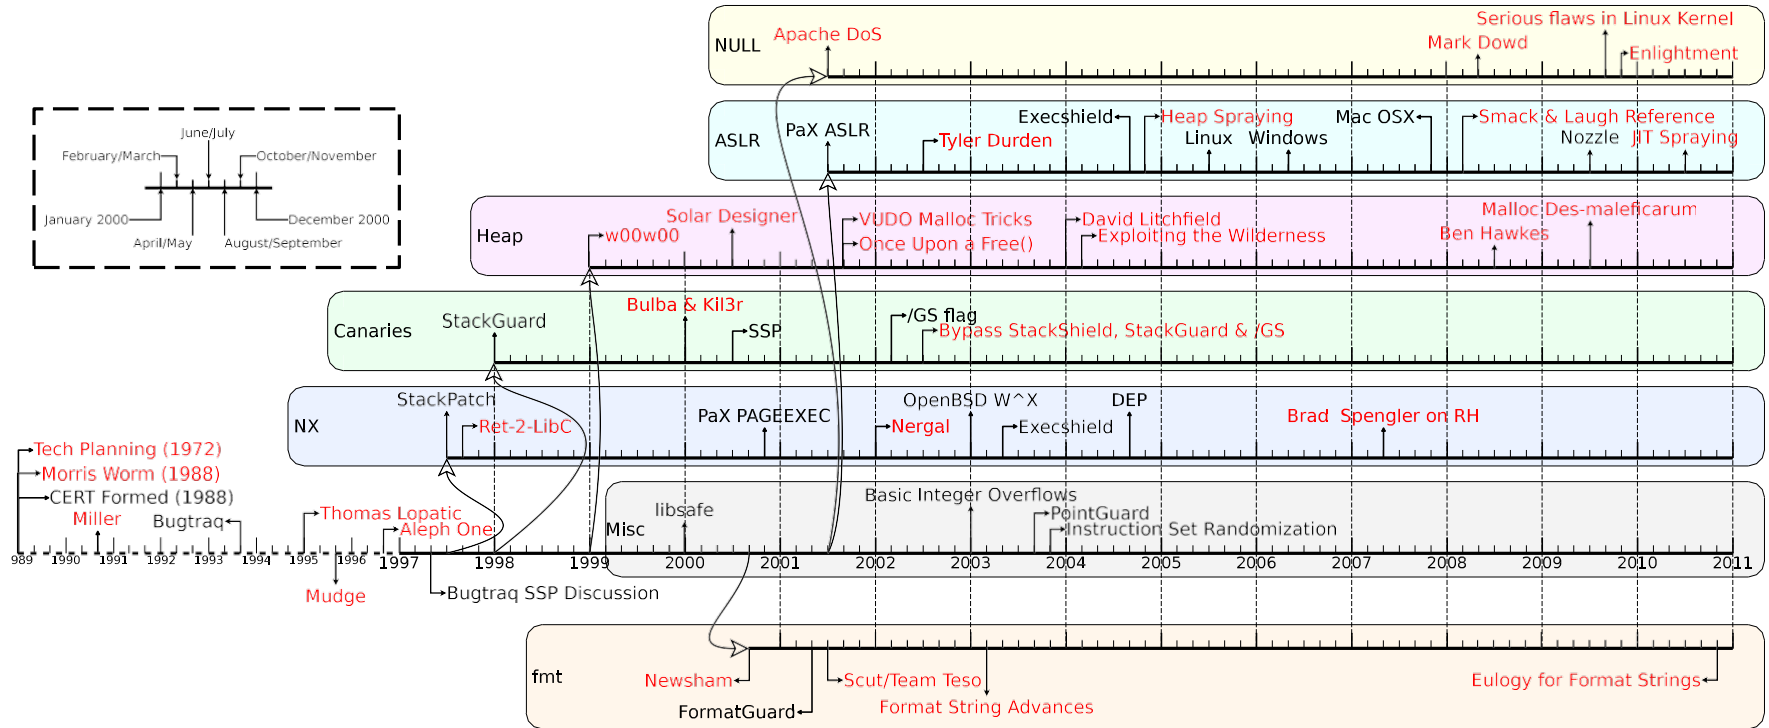
Canaries



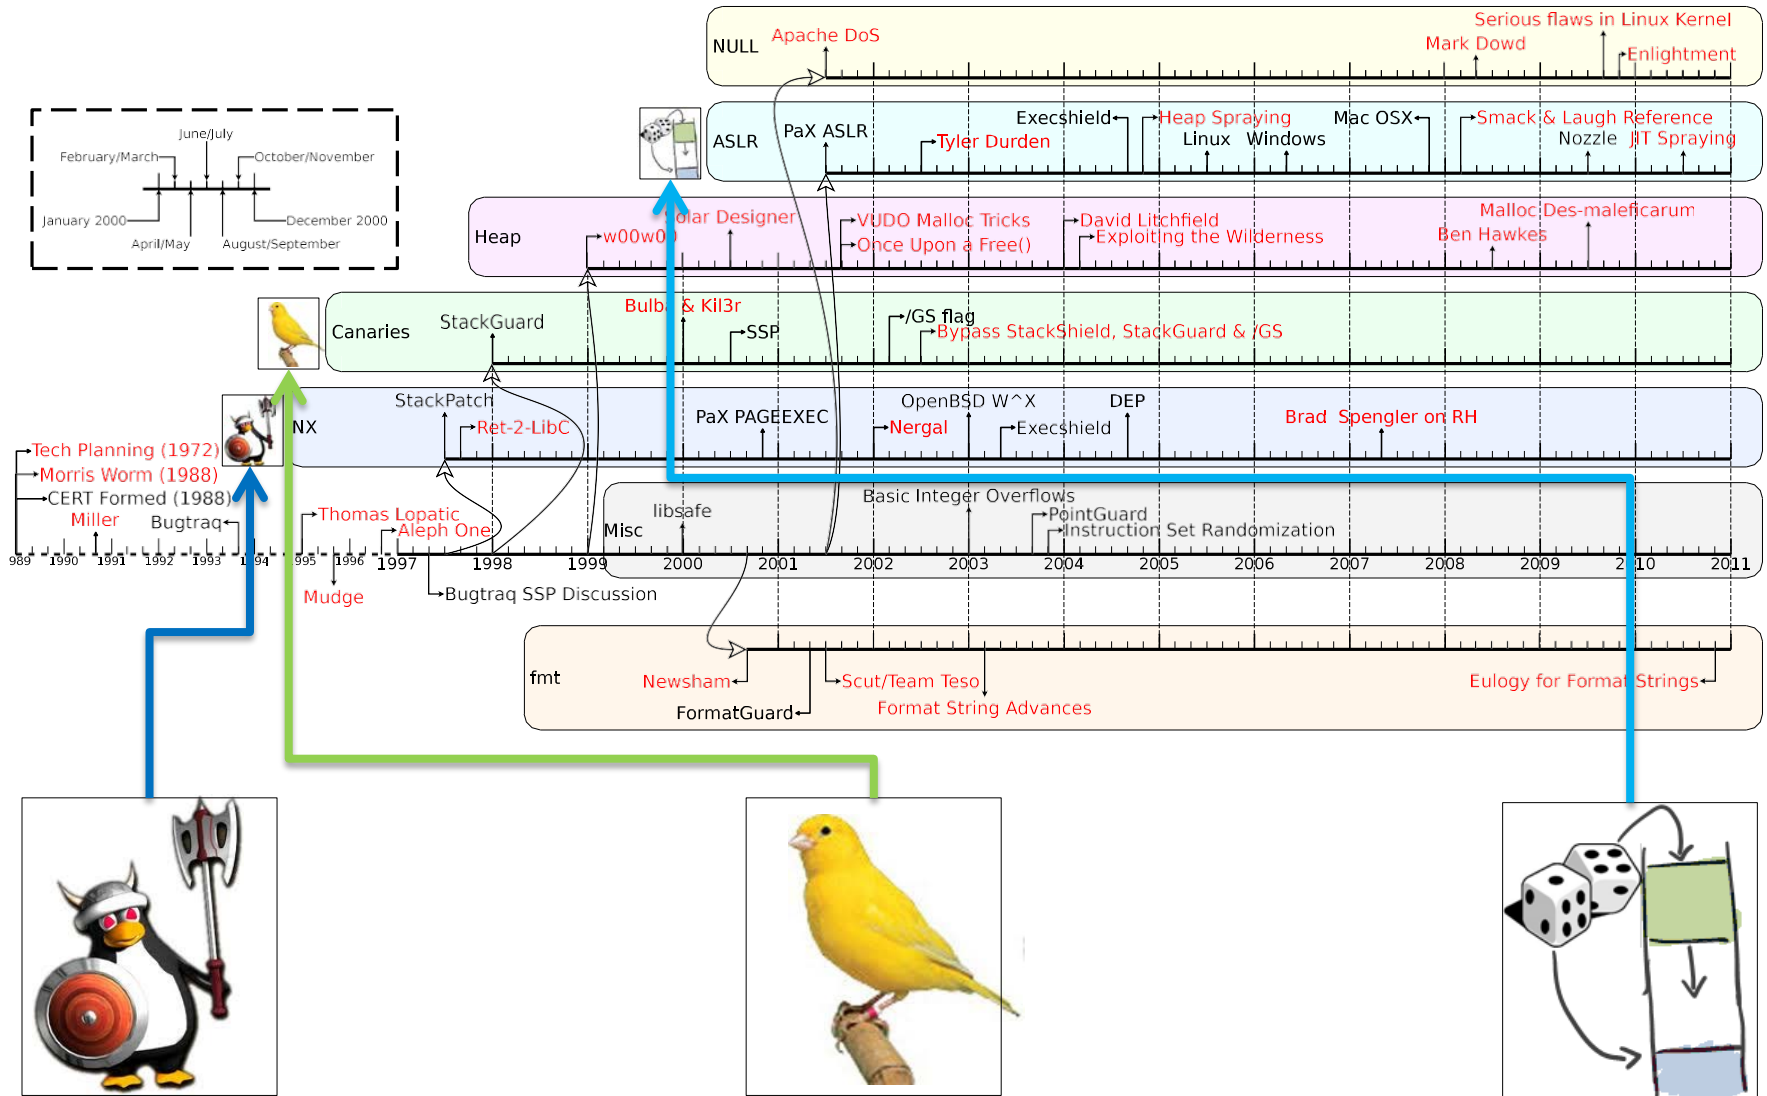
ASLR



History of memory errors



History of memory errors



III.A

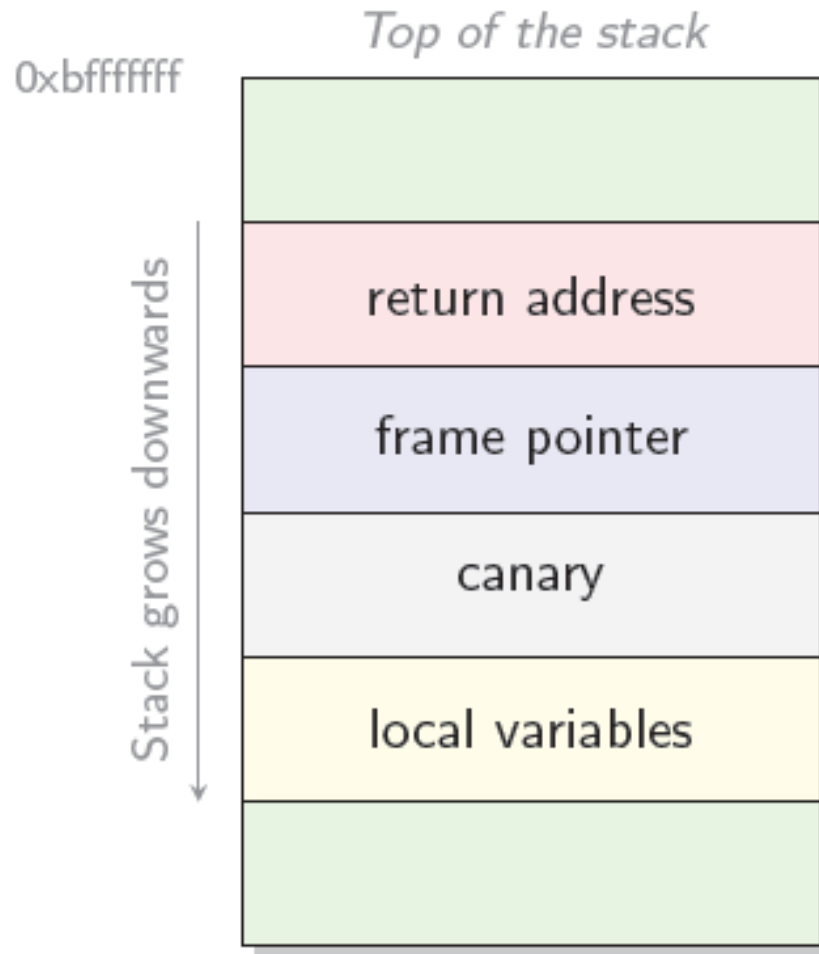
Canaries

Compiler-level techniques

Canaries

- Goal: make sure we detect overflow of return address
 - The functions' prologues insert a *canary* on the stack
 - The canary is a 32-bit value inserted between the return address and local variables
- Types of canaries:
 1. Terminator
 2. Random
 3. Random XOR
- The epilogue checks if the canary has been altered
- Drawback: requires recompilation

Canaries



How good are they?

- Assume random canaries protect the stack

Can you still exploit this?

```
char gWelcome [] = "Welcome to our system! ";
```

```
void echo (int fd)
{
```

```
    int len;
```

```
    char name [64], reply [128];
```

```
    len = strlen (gWelcome);
```

```
    memcpy (reply, gWelcome, len);
```

```
    write_to_socket (fd, "Type your name: ");
```

```
    read (fd, name, 128);
```

```
    memcpy (reply+len, name, 64);
```

```
    write (fd, reply, len + 64);
```

```
    return;
```

```
}
```

```
void server (int sockfd) {
```


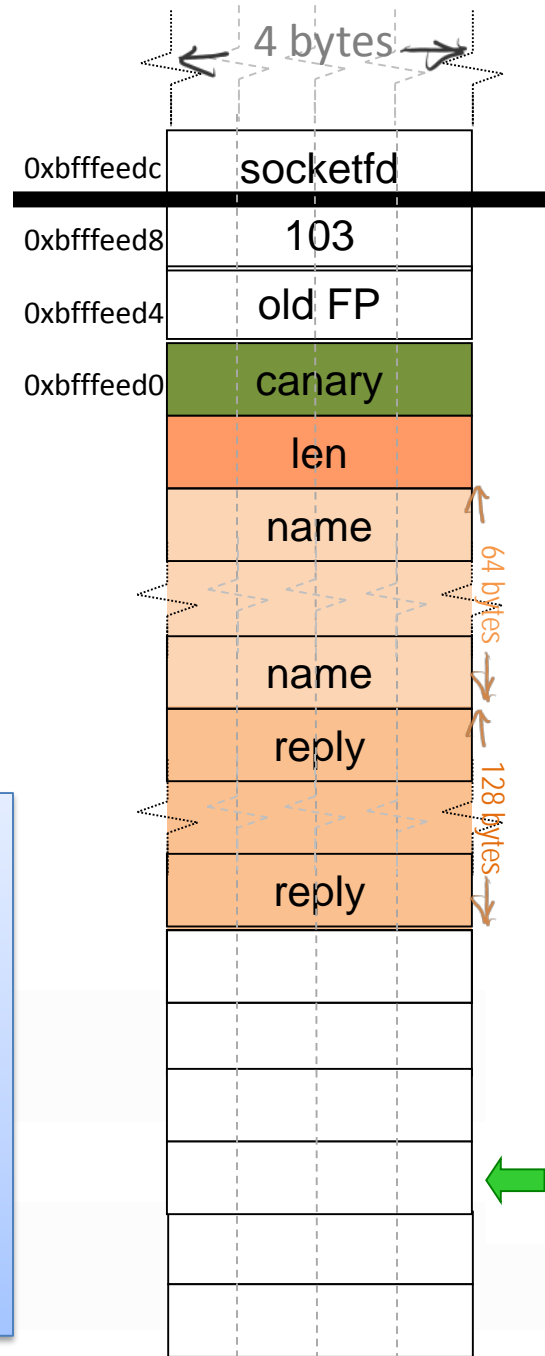
```
    while (1)
```

```
        echo (sockfd);
```

```
}
```


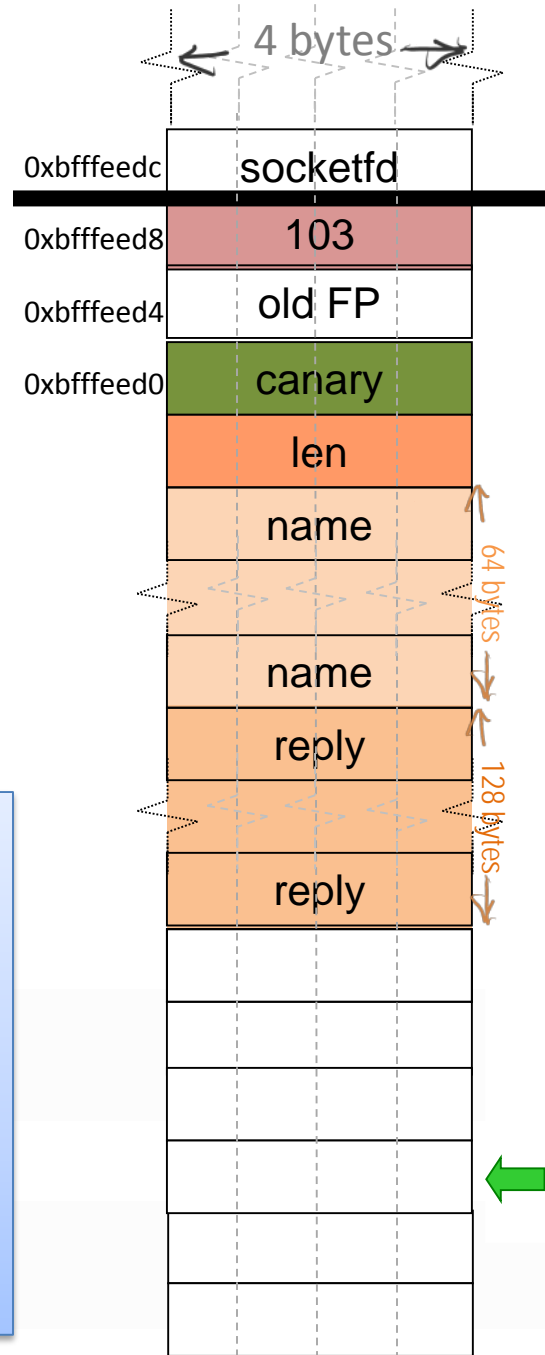

Code (echo)

```
len = strlen (gWelcome);  
memcpy (reply, gWelcome, len);  
  
read (fd, name, 128)  
  
memcpy (reply+len, name, 64)  
write (fd, reply, len +64);  
  
return;
```



Code (echo)

```
len = strlen (gWelcome);  
memcpy (reply, gWelcome, len);  
  
read (fd, name, 128)  
  
memcpy (reply+len, name, 64)  
write (fd, reply, len +64);  
  
return;
```




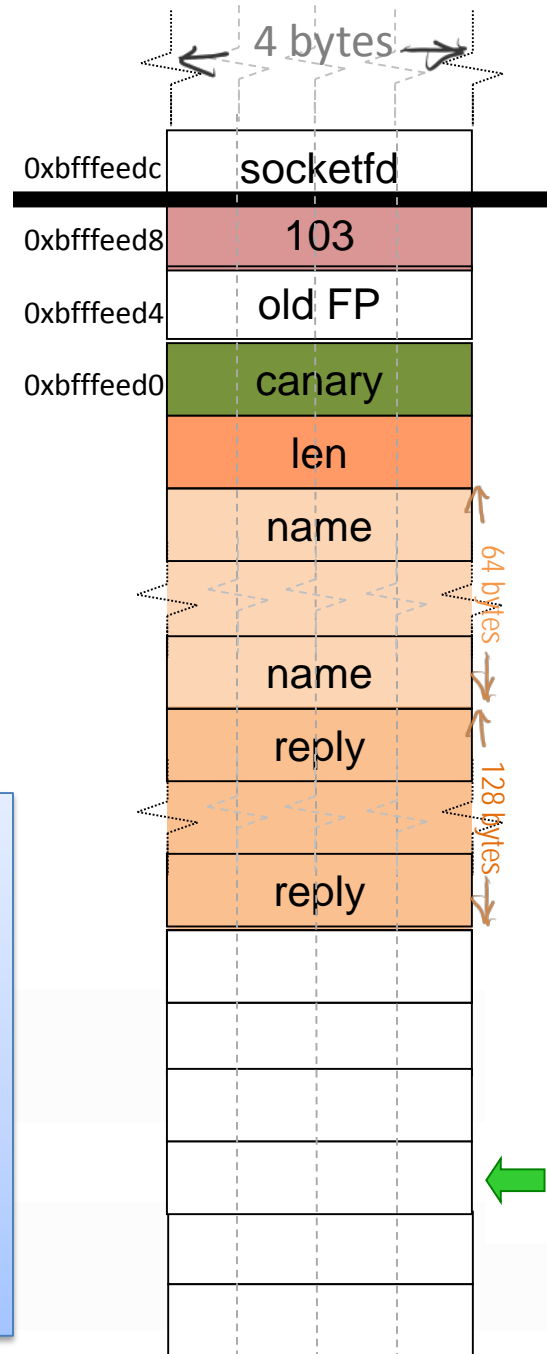
Code (echo)

```
len = strlen (gWelcome);  
memcpy (reply, gWelcome, len);
```

```
read (fd, name, 128) ←
```

```
memcpy (reply+len, name, 64)  
write (fd, reply, len +64);
```

```
return;
```



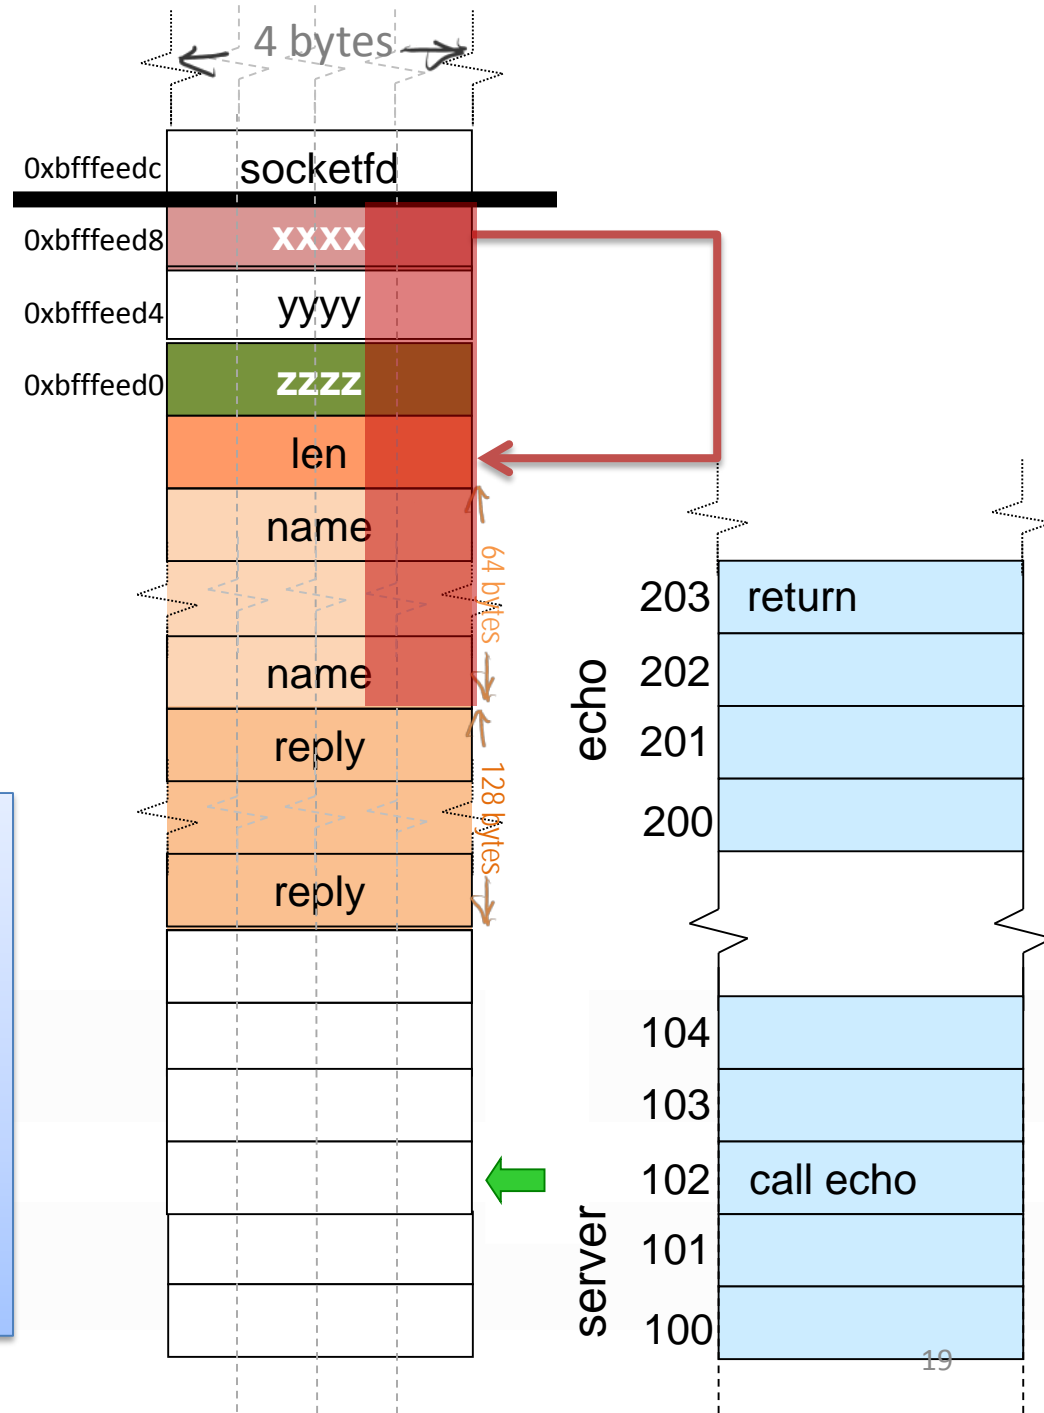
Code (echo)

```
len = strlen (gWelcome);  
memcpy (reply, gWelcome, len);
```

```
read (fd, name, 128) ←
```

```
memcpy (reply+len, name, 64)  
write (fd, reply, len +64);
```

```
return;
```



Will not be the same value, so program will crash. OK for causing problems but we can no longer run our own code.

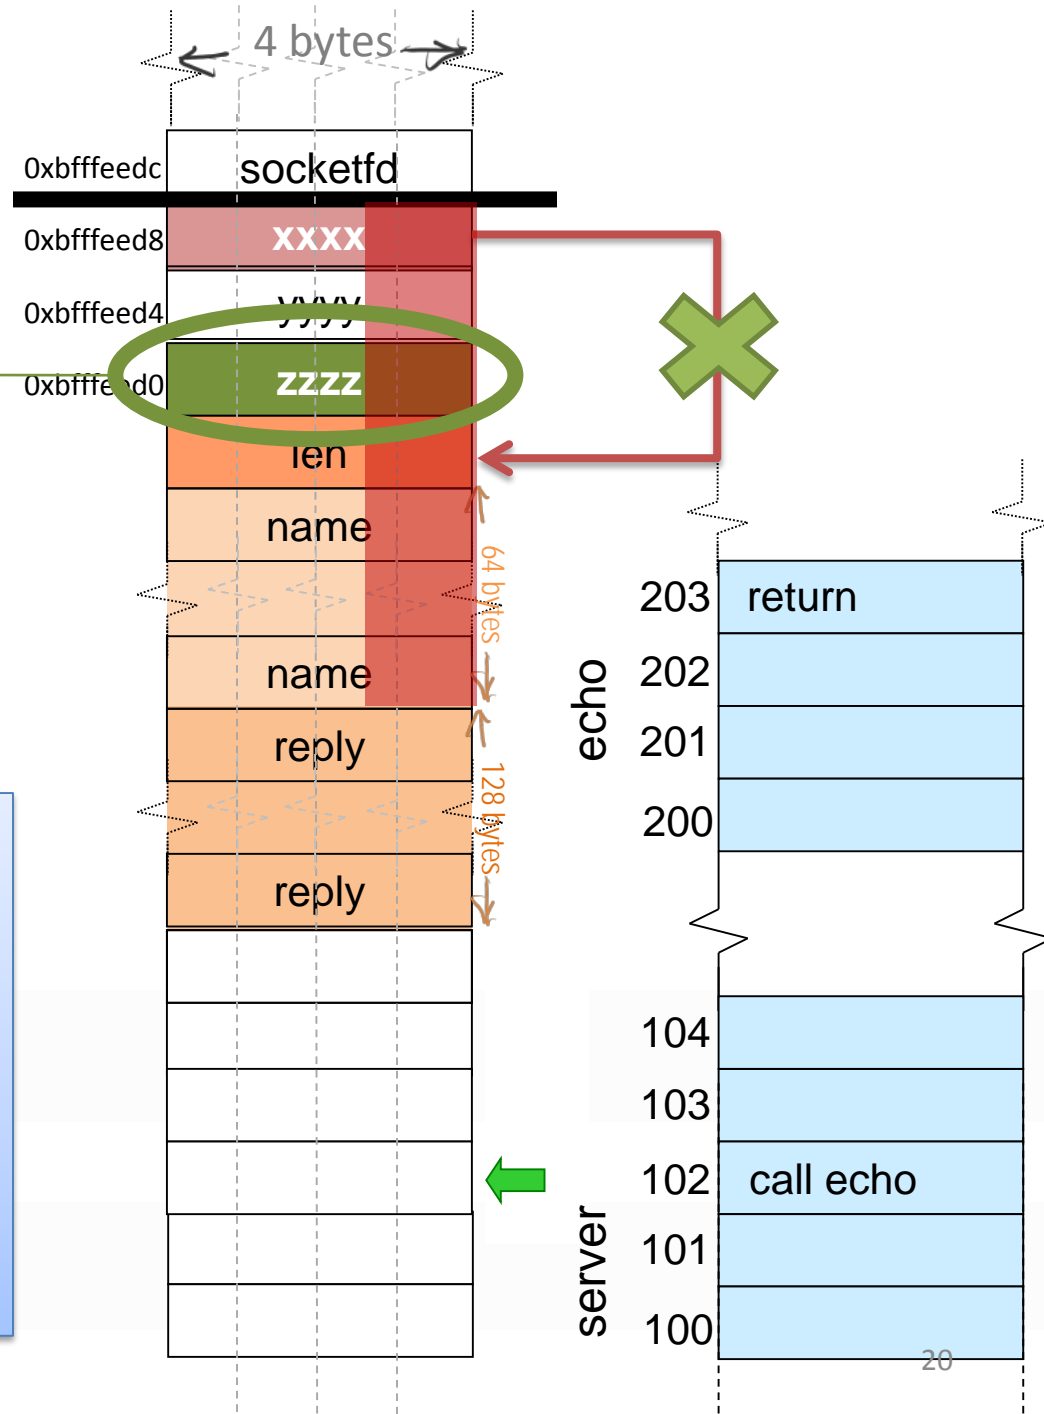
Code (echo)

```
len = strlen(gWelcome);
memcpy(reply, gWelcome, len);
```

read (fd, name, 128) ←

```
memcpy(reply+len, name, 64)
write (fd, reply, len +64);
```

```
return;
```


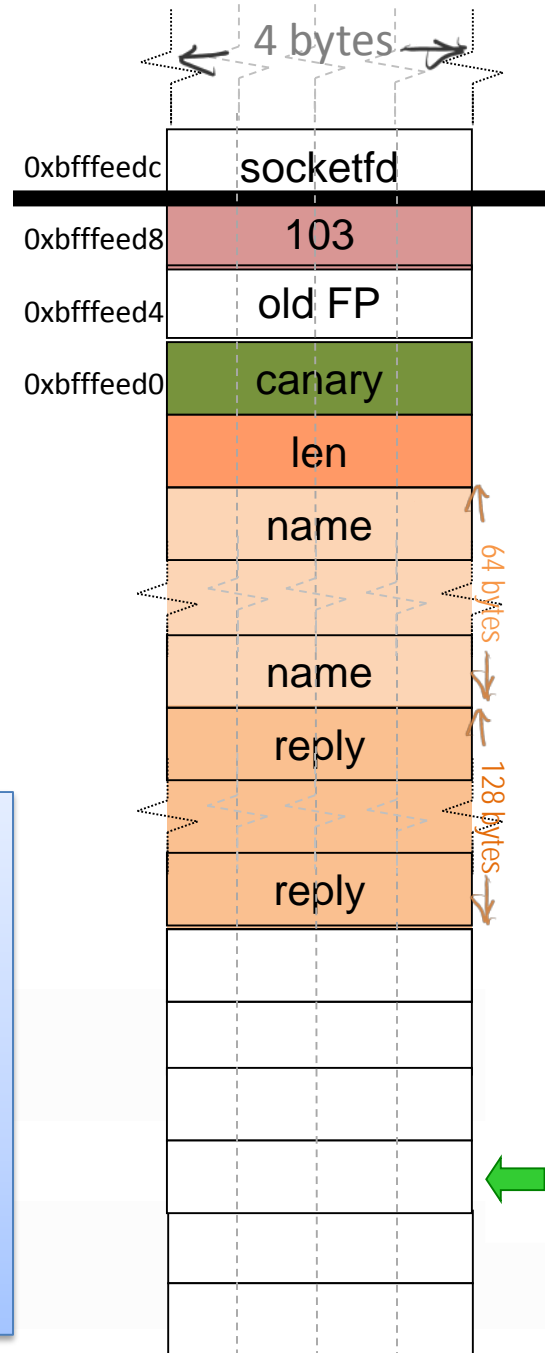


Are we safe?

Any ideas?


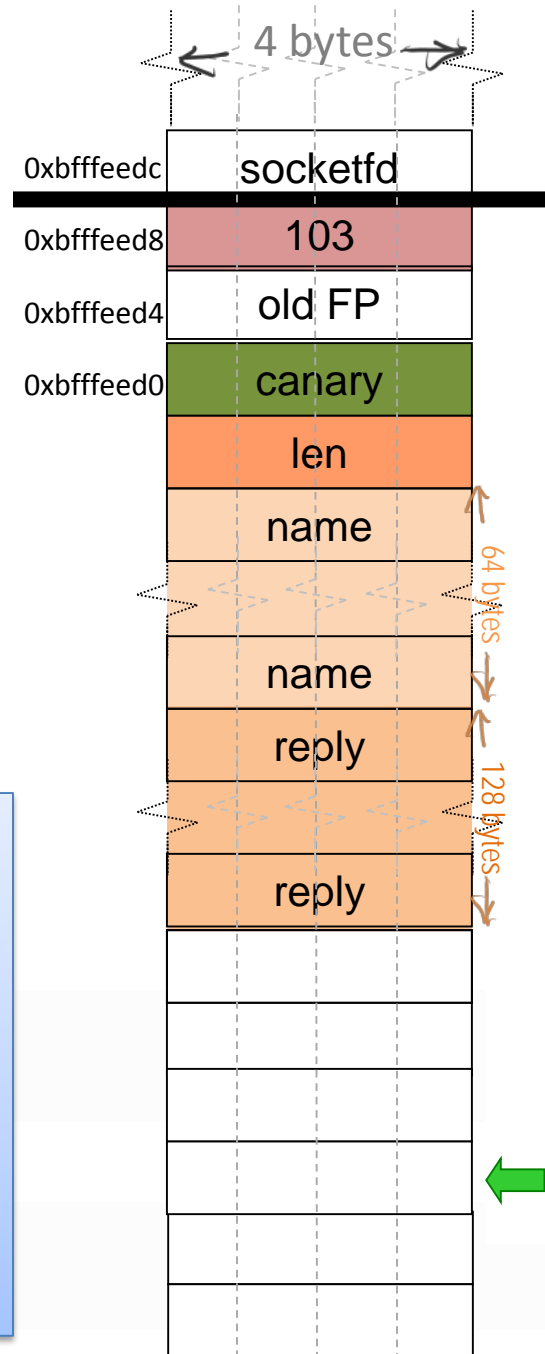
Code (echo)

```
len = strlen (gWelcome);  
memcpy (reply, gWelcome, len);  
  
read (fd, name, 128)  
  
memcpy (reply+len, name, 64)  
write (fd, reply, len +64);  
  
return;
```



Code (echo)

```
len = strlen(gWelcome);  
memcpy(reply, gWelcome, len);  
  
read(fd, name, 128)  
  
memcpy(reply+len, name, 64) ←  
write(fd, reply, len + 64);  
  
return;
```





Code (echo)

```
len = strlen (gWelcome);  
memcpy (reply, gWelcome, len);
```

```
read (fd, name, 128) ←
```

```
memcpy (reply+len, name, 64)  
write (fd, reply, len +64);
```

```
return;
```





Code (echo)

```
len = strlen (gWelcome);  
memcpy (reply, gWelcome, len);
```

```
read (fd, name, 128) ←
```

```
memcpy (reply+len, name, 64)  
write (fd, reply, len +64);
```

```
return;
```



Will **STILL** pass we
can run our own
code.

Code (echo)

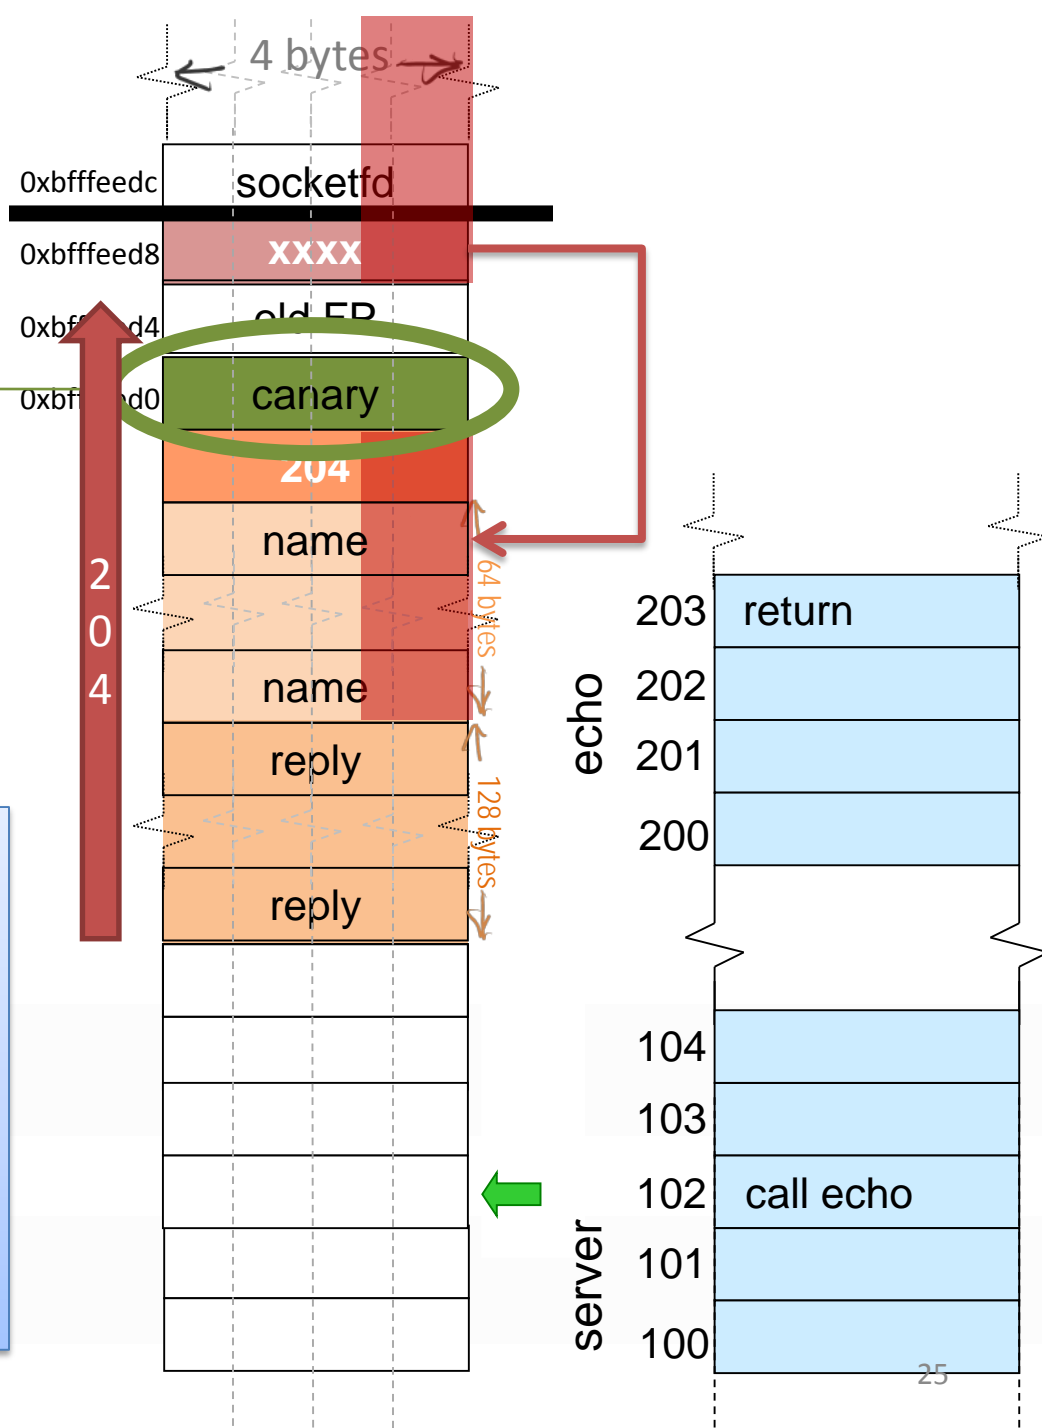
```
len = strlen (gWelcome);  
memcpy (reply, gWelcome, len);
```

```
read (fd, name, 128)
```

```
memcpy (reply+len, name, 64)
```

```
write (fd, reply, len +64);
```

```
return;
```



III.B

“DEP”

DEP / NX bit / W \oplus X

- Idea: separate executable memory locations from writable ones
 - A memory page cannot be both writable and executable at the same time
- “Data Execution Prevention (DEP)”

Code (echo)

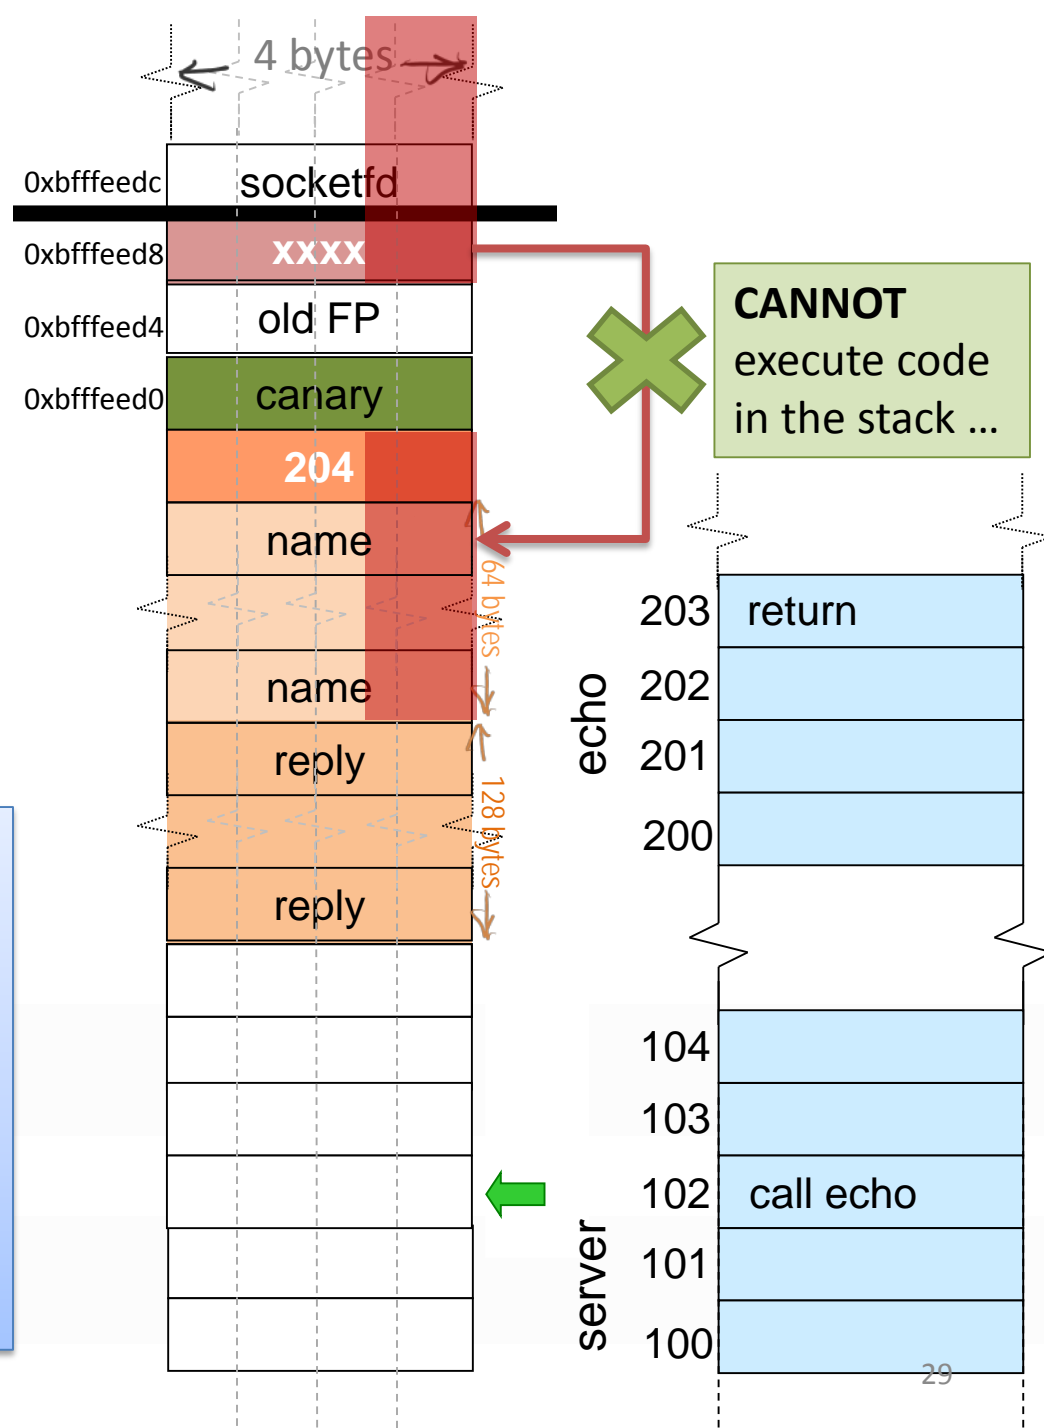
```
len = strlen (gWelcome);  
memcpy (reply, gWelcome, len);
```

```
read (fd, name, 128)
```

```
memcpy (reply+len, name, 64) ←
```

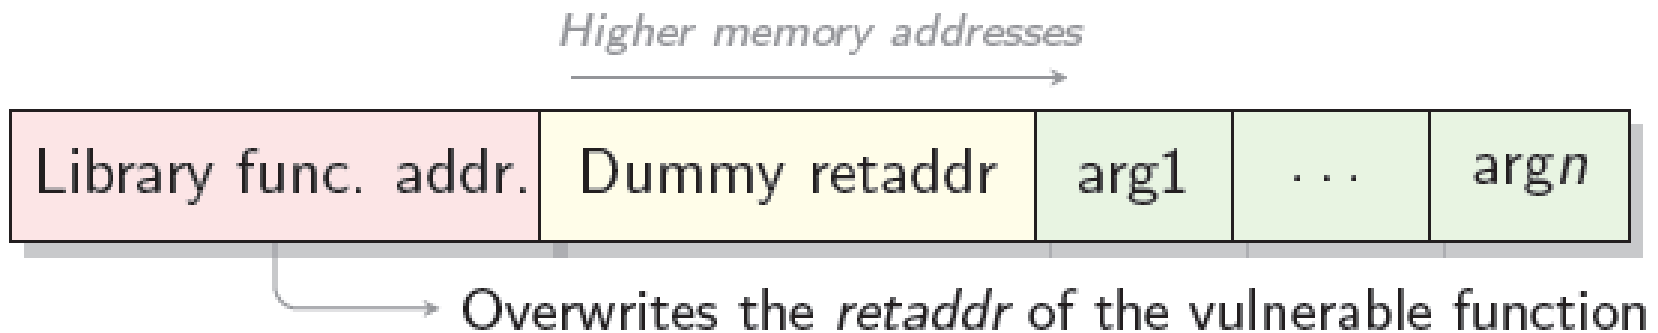
```
write (fd, reply, len +64);
```

```
return;
```

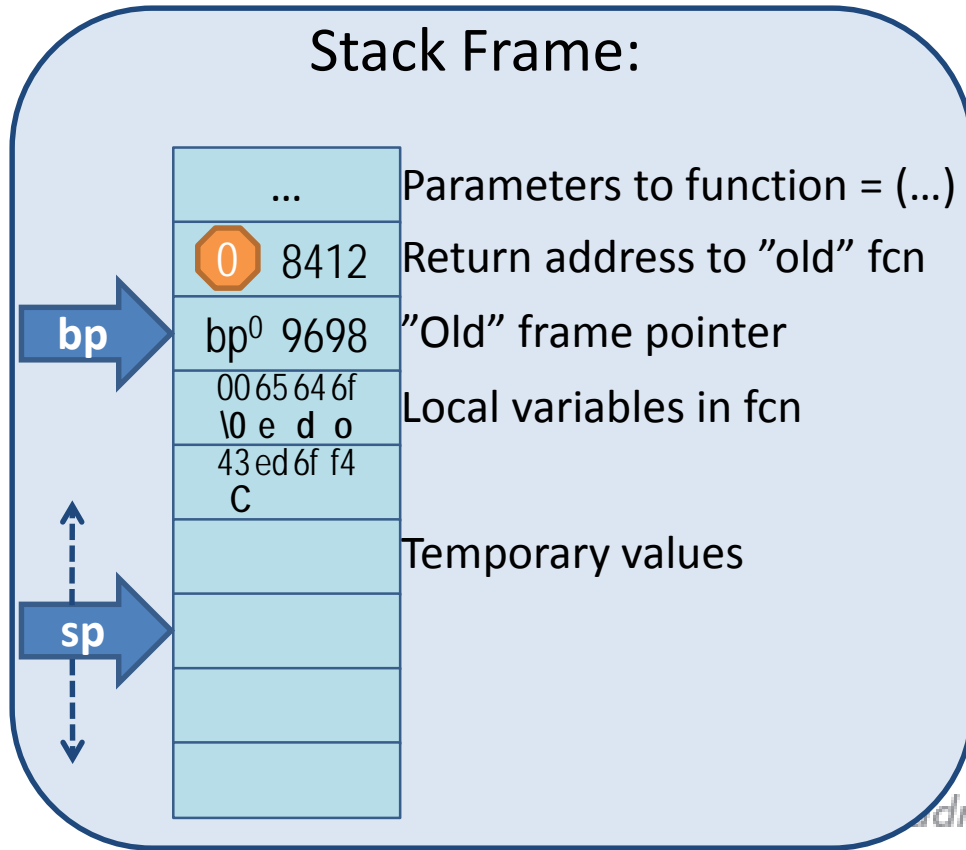


Bypassing $W \oplus X$

- Return into libc
- Three assumptions:
 - We can manipulate a code pointer
 - The stack is writable
 - We know the address of a “suitable” library function (e.g., `system()`)

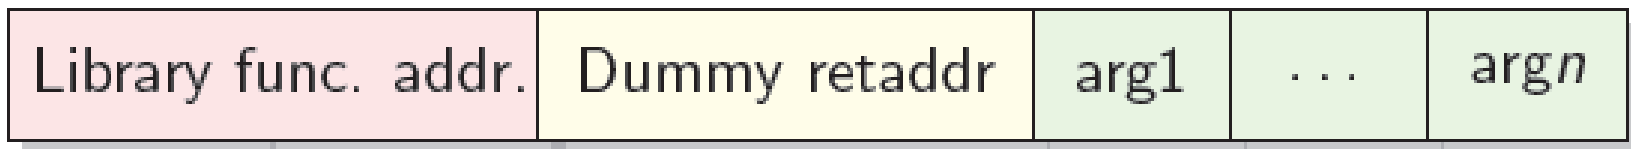


Bypassing $W \oplus X$




When **calling** a function:

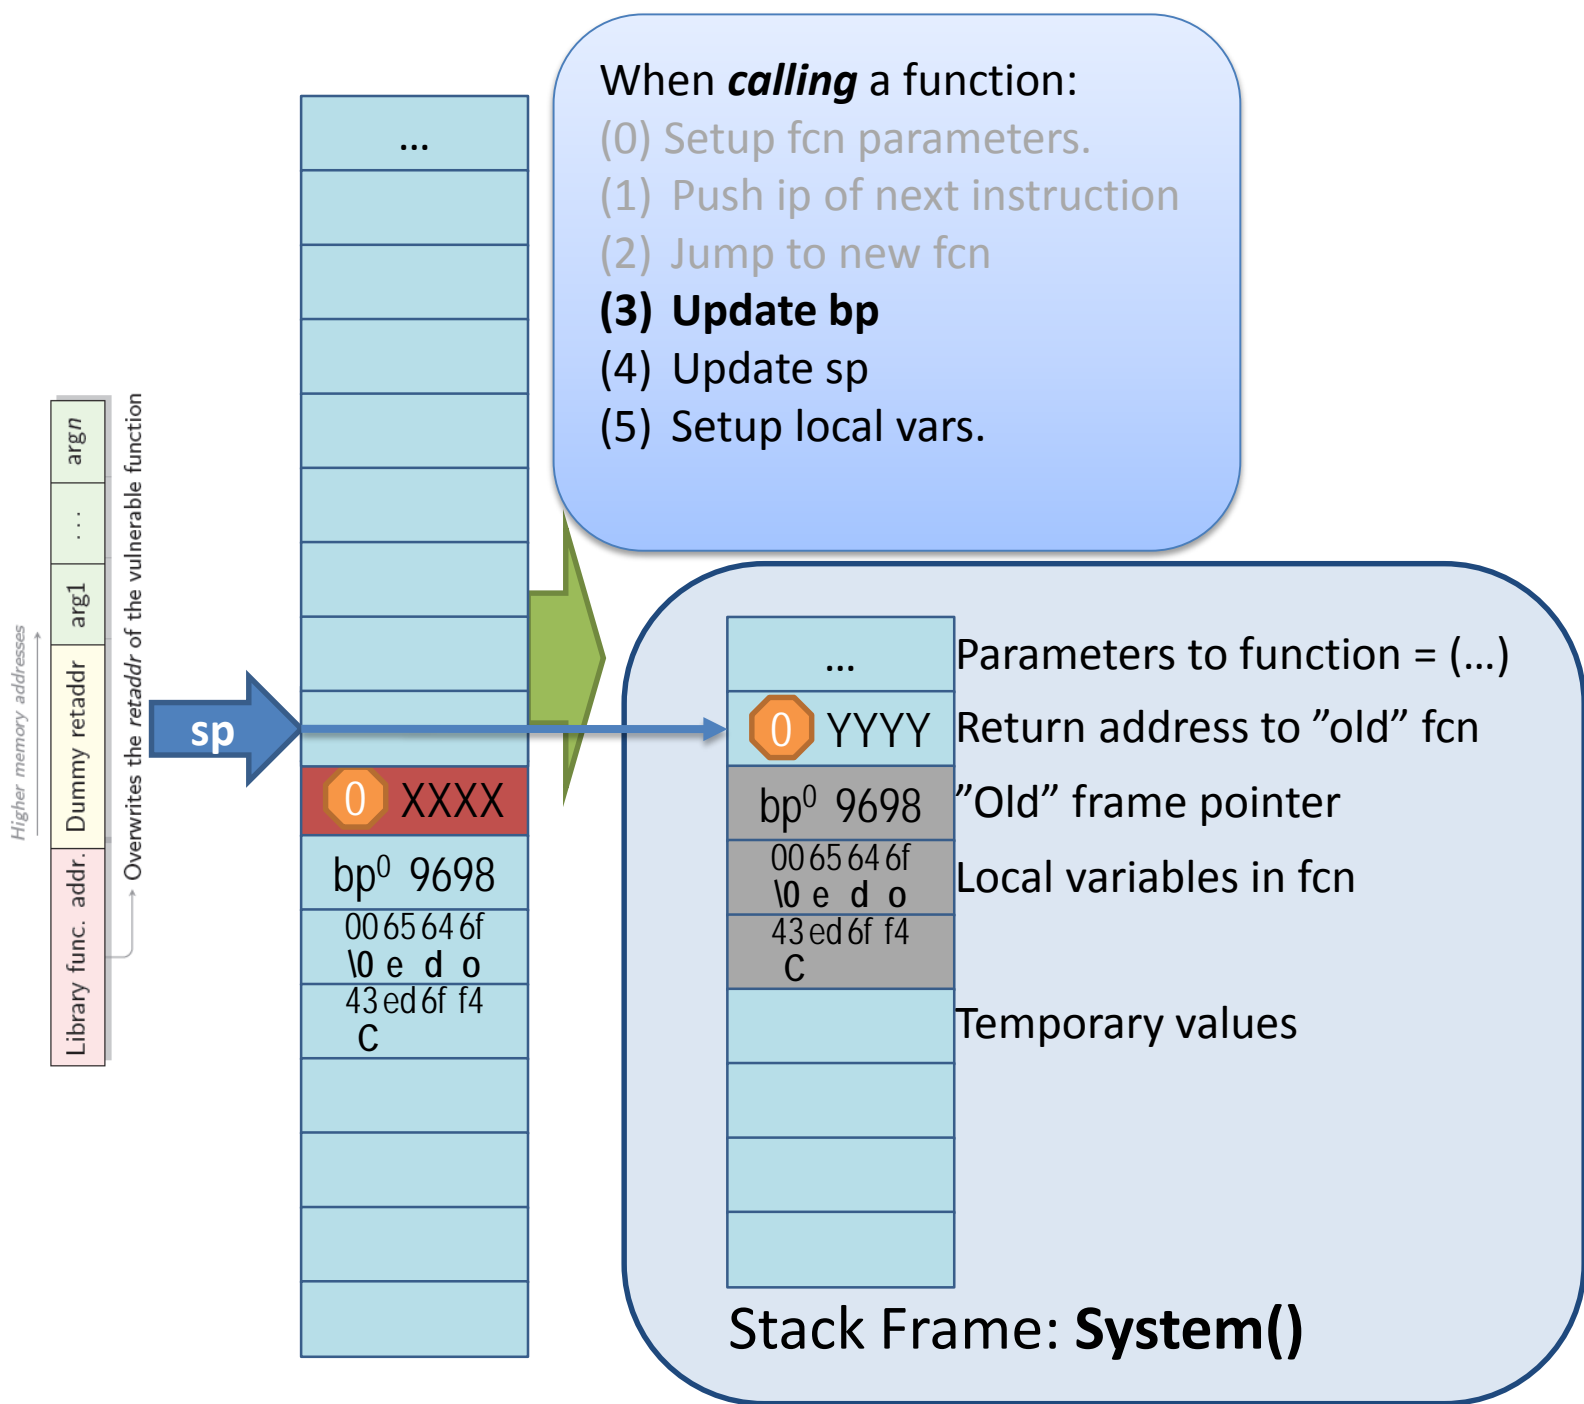
- (0) Setup fcn parameters.
- (1) Push ip of next instruction
- (2) Jump to new fcn
- (3) Update bp**
- (4) Update sp
- (5) Setup local vars.



Overwrites the *retaddr* of the vulnerable function





Jump to the "fcn"
we want to execute,
for example:
system()





When **calling** a function:

- (0) Setup fcn parameters.
- (1) Push ip of next instruction
- (2) Jump to new fcn
- (3) Update bp**
- (4) Update sp
- (5) Setup local vars.



Parameters to function = (...)

Return address to "old" fcn

"Old" frame pointer


Local variables in fcn

Temporary values

Stack Frame: **System()**




sp



When **calling** a function:

- (0) Setup fcn parameters.
- (1) Push ip of next instruction
- (2) Jump to new fcn
- (3) Update bp**
- (4) Update sp
- (5) Setup local vars.

Why the “ret address”?
What could we do with it?



Parameters to function = (...)

Return address to “old” fcn

“Old” frame pointer


Local variables in fcn

Temporary values

Stack Frame: **System()**


Stack

- Why the “ret address”?
- What could we do with it?



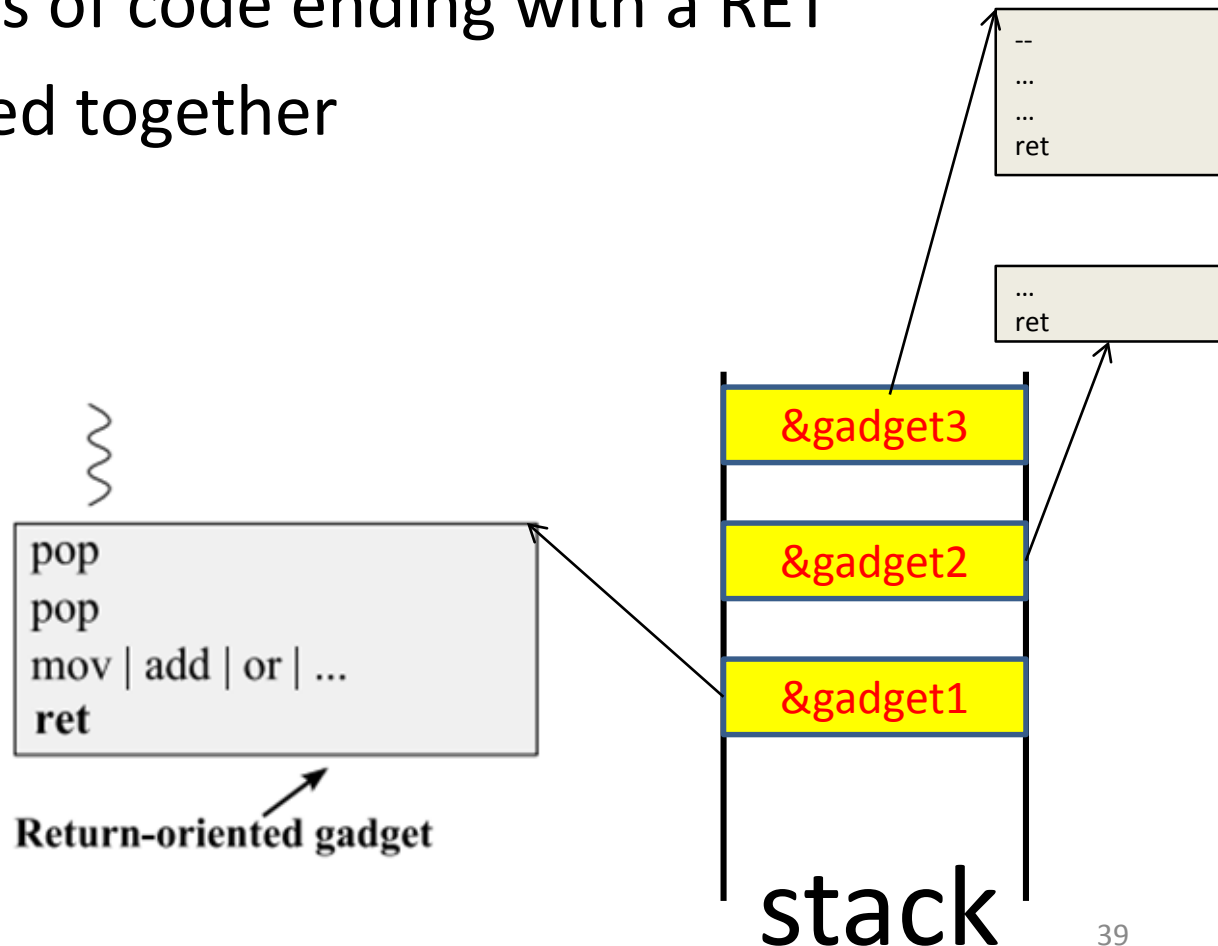
Return Oriented Programming

- ROP chains:
 - Small snippets of code ending with a RET
 - Can be chained together



Return Oriented Programming

- ROP chains
 - Small snippets of code ending with a RET
 - Can be chained together



How good are they?

- Assume random canaries protect the stack
- Assume DEP prevents execution of the stack

Can you still exploit this?

```
char gWelcome [] = "Welcome to our system! ";
```

```
void echo (int fd)
```

```
{
```

```
    int len;
```

```
    char name [64], reply [128];
```

```
    len = strlen (gWelcome);
```

```
    memcpy (reply, gWelcome, len);
```

```
    write_to_socket (fd, "Type your name: ");
```

```
    read (fd, name, 128);
```

```
    memcpy (reply+len, name, 64);
```

```
    write (fd, reply, len + 64);
```

```
    return;
```

```
}
```

```
void server (int sockfd) {
```

```
    while (1)
```

```
        echo (sockfd);
```

```
}
```

III.C

ASLR

Let us make it a little harder still...

Address Space Layout Randomisation

- Idea:
 - Re-arrange the position of key data areas randomly (stack, .data, .text, shared libraries, . . .)
 - Buffer overflow: the attacker does not know the address of the shellcode
 - Return-into-libc: the attacker can't predict the address of the library function
 - Implementations: Linux kernel > 2.6.11, Windows Vista, . . .

ASLR: Problems

- 32-bit implementations use few randomisation bits
- An attacker can still exploit non-randomised areas, or rely on other information leaks (e.g., format bug)
- So... (I bet you saw this one coming)....

How good are they?

- Assume random canaries protect the stack
- Assume DEP prevents execution of the stack
- Assume ASLR randomized the stack *and* the start address of the code
 - but let us assume that all functions are still at the same relative offset from start address of code
 - (in other words: need only a single code pointer)

Can you still exploit this?

```
char gWelcome [] = "Welcome to our system! ";
```

```
void echo (int fd)
```

```
{
```

```
    int len;
```

```
    char name [64], reply [128];
```

```
    len = strlen (gWelcome);
```

```
    memcpy (reply, gWelcome, len);
```

```
    write_to_socket (fd, "Type your name: ");
```

```
    read (fd, name, 128);
```

```
    memcpy (reply+len, name, 64);
```

```
    write (fd, reply, len + 64);
```

```
    return;
```

```
}
```

```
void server (int sockfd) {
```

```
    while (1)
```

```
        echo (sockfd);
```

```
}
```


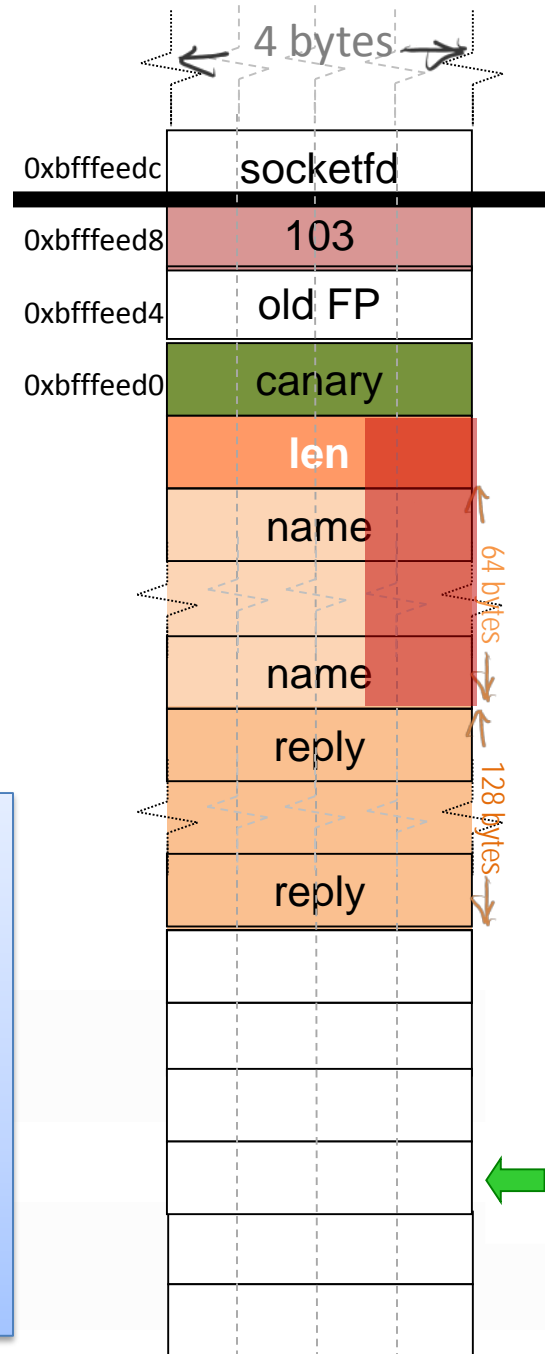
Code (echo)

```
len = strlen(gWelcome);
memcpy(reply, gWelcome, len);
```

read (fd, name, 128) ←

```
memcpy(reply+len, name, 64)
write (fd, reply, len +64);
```

```
return;
```



Code (echo)


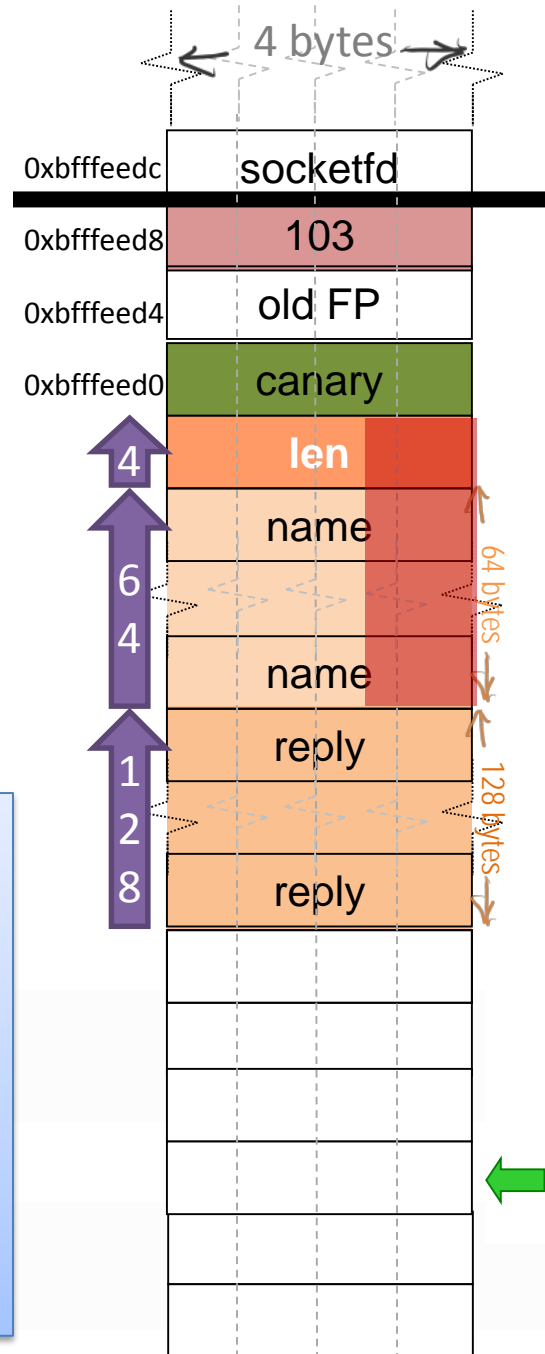
```
len = strlen(gWelcome);
memcpy(reply, gWelcome, len);
```

```
read(fd, name, 128)
```

```
memcpy(reply+len, name, 64) ←
```

```
write(fd, reply, len + 64);
```

```
return;
```



Code (echo)


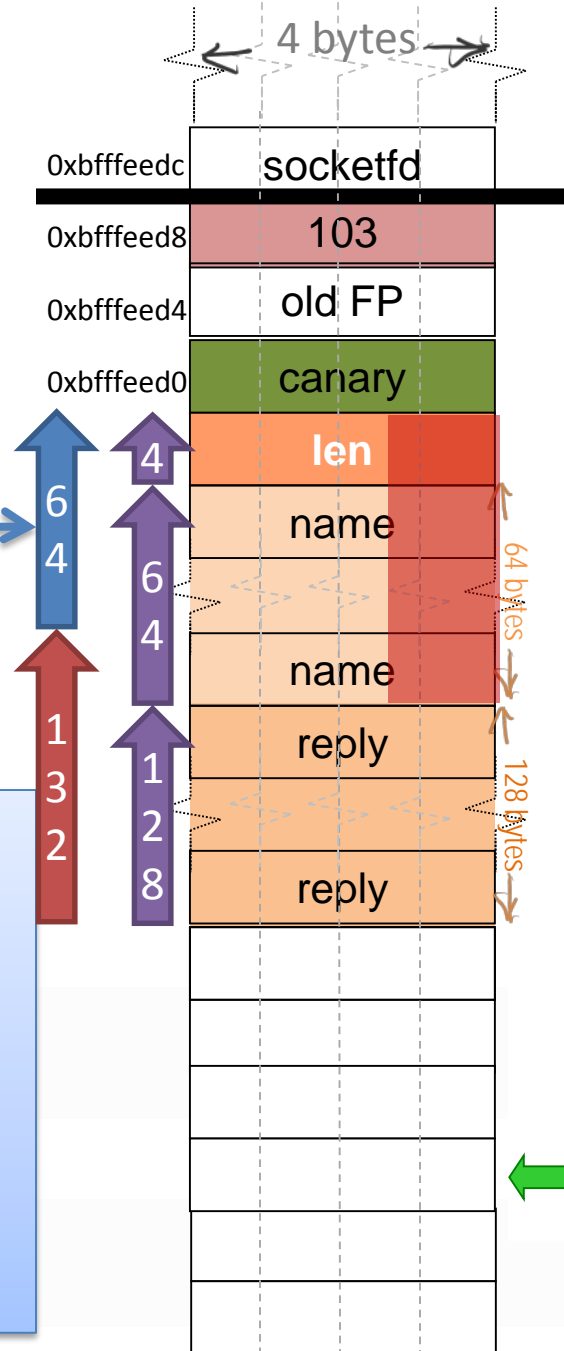
```
len = strlen (gWelcome);
memcpy (reply, gWelcome, len);
```

```
read (fd, name, 128)
```

```
memcpy (reply+len, name, 64) ←
```

```
write (fd, reply, len +64);
```

```
return;
```



Code (echo)


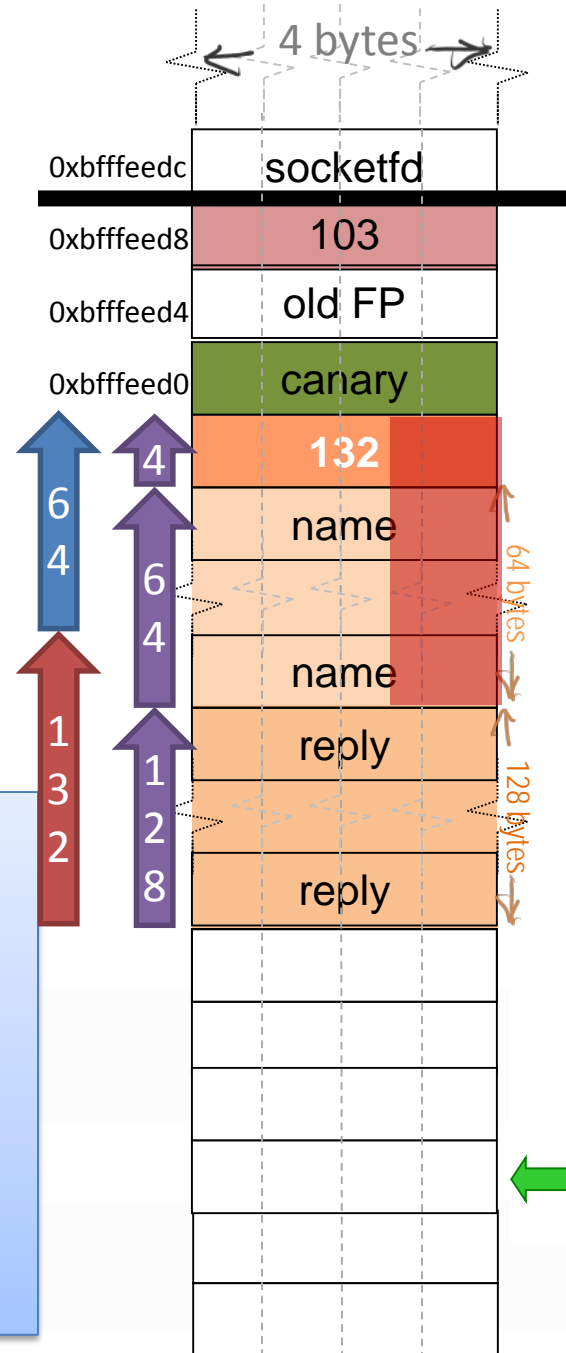
```
len = strlen (gWelcome);
memcpy (reply, gWelcome, len);
```

```
read (fd, name, 128)
```

```
memcpy (reply+len, name, 64) ←
```

```
write (fd, reply, len +64);
```

```
return;
```



Code (echo)


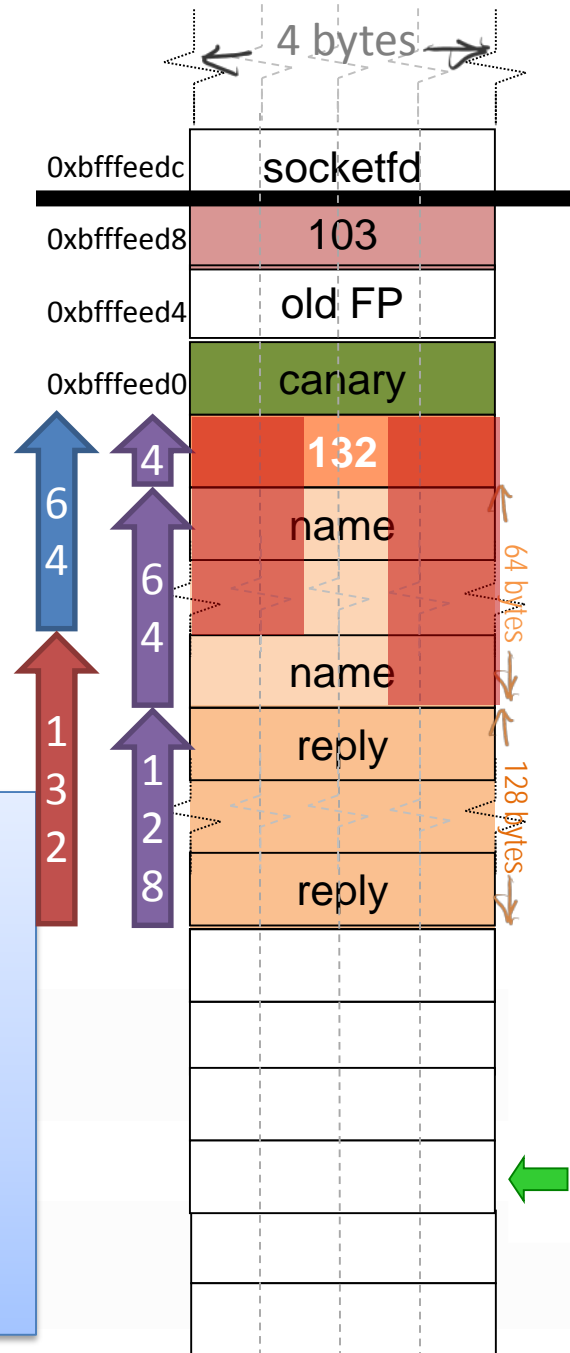
```
len = strlen (gWelcome);
memcpy (reply, gWelcome, len);
```

```
read (fd, name, 128)
```

```
memcpy (reply+len, name, 64) ←
```

```
write (fd, reply, len +64);
```

```
return;
```



Code (echo)


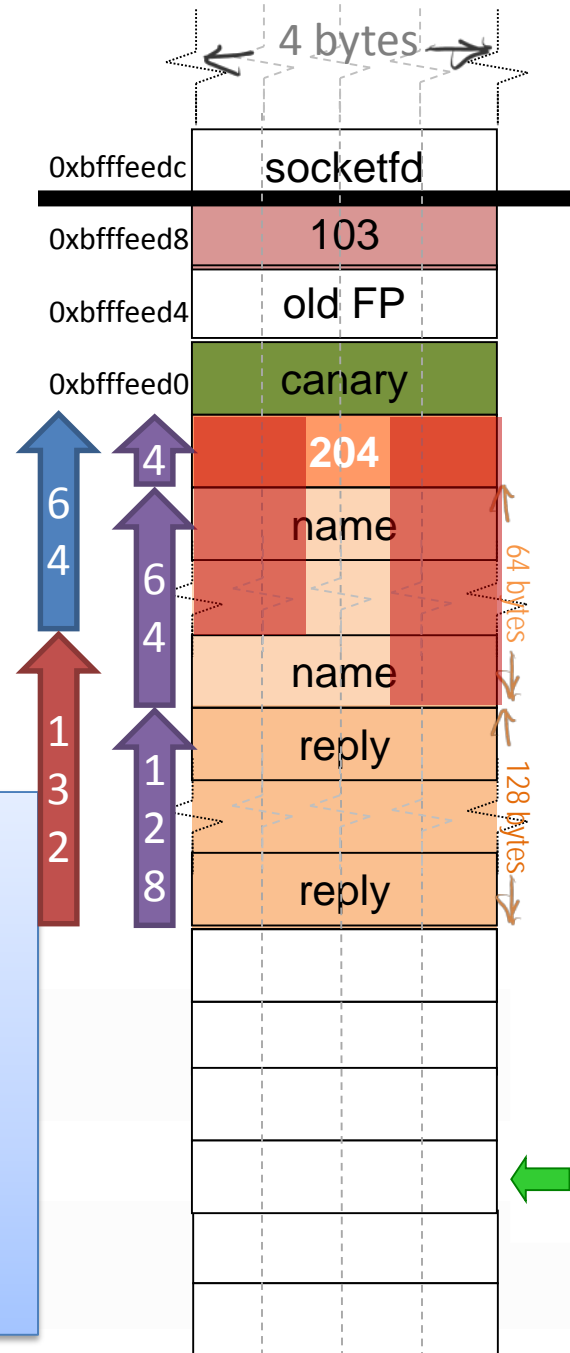
```
len = strlen (gWelcome);
memcpy (reply, gWelcome, len);
```

```
read (fd, name, 128)
```

```
memcpy (reply+len, name, 64) ←
```

```
write (fd, reply, len +64);
```

```
return;
```



Code (echo)


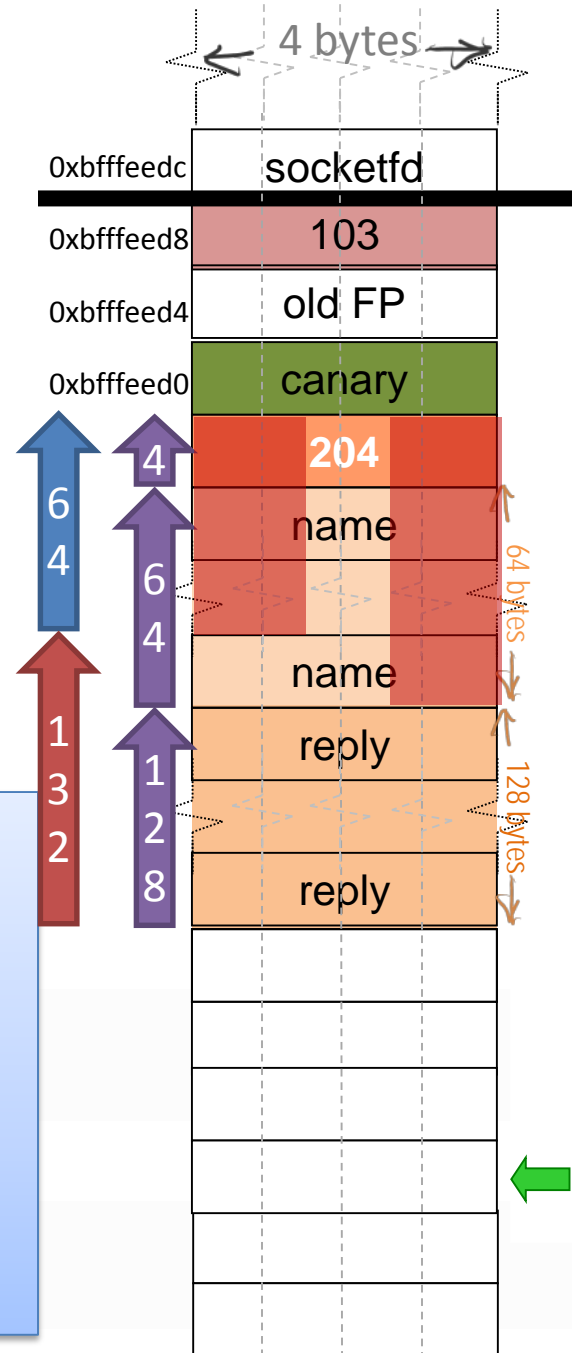
```
len = strlen (gWelcome);
memcpy (reply, gWelcome, len);
```

```
read (fd, name, 128)
```

```
memcpy (reply+len, name, 64)
```

```
write (fd, reply, len +64);
```

```
return;
```



Code (echo)


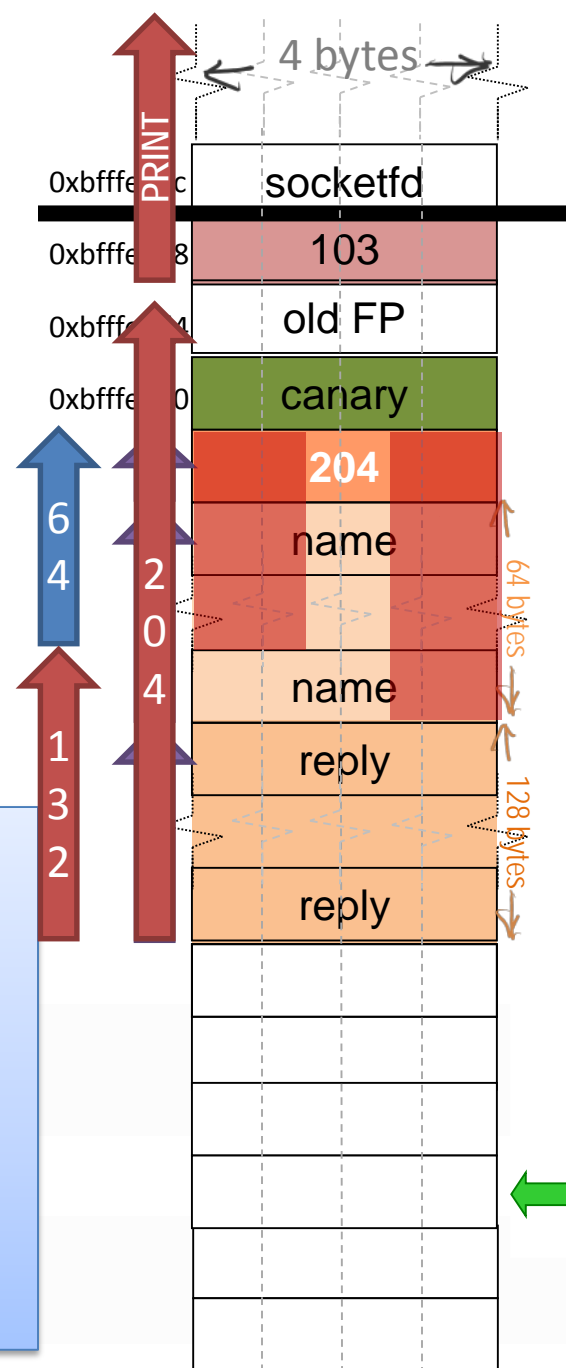
```
len = strlen (gWelcome);
memcpy (reply, gWelcome, len);
```

```
read (fd, name, 128)
```

```
memcpy (reply+len, name, 64)
```

```
write (fd, reply, len +64);
```

```
return;
```



Print a code address.

Then rerun program
because then we can
defeat ASLR.

Code (echo)


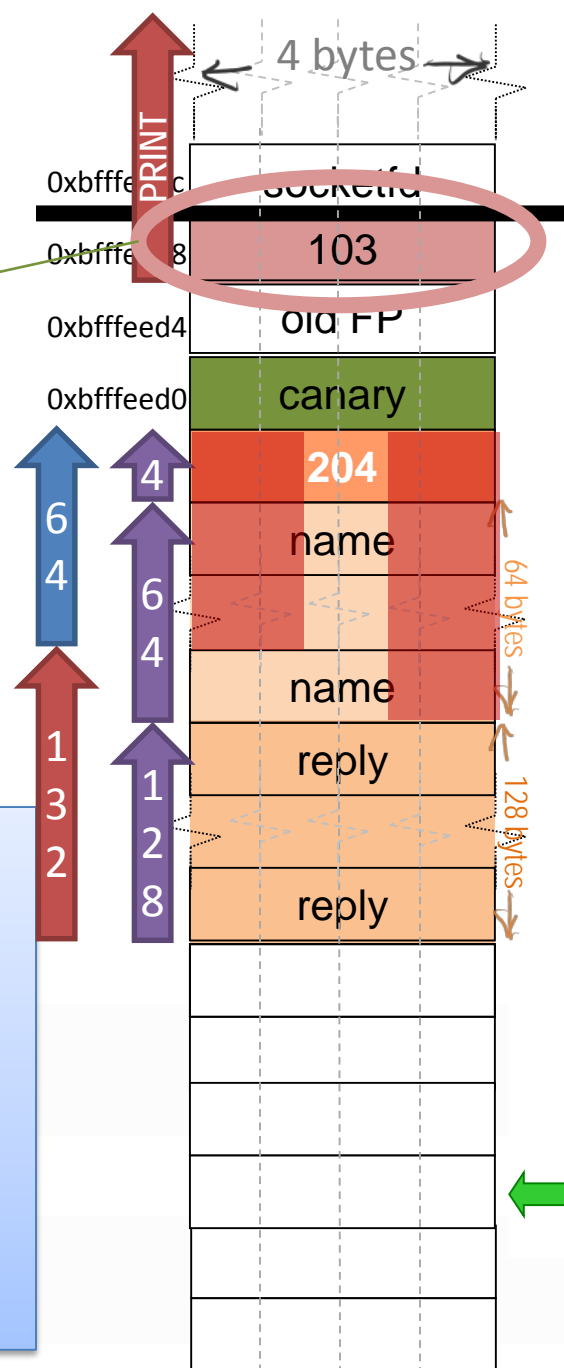
```
len = strlen (gWelcome);  
memcpy (reply, gWelcome, len);
```

```
read (fd, name, 128)
```

```
memcpy (reply+len, name, 64)
```

```
write (fd, reply, len +64);
```

```
return;
```




Finally

You may also overwrite other things

- For instance:
 - Other variables that are also on the stack
 - Other addresses
 - Etc.

Exploit against non-control data

```
get_medical_info()  
{  
    boolean authorized = false;  
    char name [10];  
    authorized = check();  
    read_from_network (name);  
  
    if (authorized)  
        show_medical_info (name);  
    else  
        printf ("sorry, not allowed");  
}
```



Memory Corruption

Summary

- We have sketched only the most common memory corruption attack
 - many variations, e.g.:
 - heap \leftrightarrow stack
 - more complex overflows
 - off-by-one
- But there are others also
 - integer overflows
 - format string attacks
 - double free
 - etc.
- Not now, perhaps later...

We constructed “weird machines”

- New spin on fundamental questions:
 - ➔ “What is computable?”
- Shellcode, ROP, Ret2Libc
 - ➔ Turing Complete





That is all folks!

- We have covered quite a lot:
 - Simple buffer overflows
 - Counter measures
 - Counter counter measures
- Research suggests that buffer overflows will be with us for quite some time
- Best avoid them in your code!